

Impact markets may incentivize predictably net-negative projects

113

by Ofer, Owen Cotton-Barratt Jun 21 2022



Building effective altruism

Certificate of impact

Accidental harm

Markets for altruism

Show all topics

Summary

[Impact markets](#)^o (that encourage retrospective funding, and especially if they allow resale of impact) have a severe downside risk: they can incentivize risky projects that are likely to be net-negative due to allowing people to profit if they cause positive impact while not inflicting a cost on them if they cause negative impact. This risk is hard to mitigate.

Impact markets themselves are therefore such a risky project. To avoid the conflict of interest issues that arise, work to establish impact markets should only ever be funded prospectively (never retrospectively).

The risk

Suppose the certificates of a risky project are traded on an impact market. If the project ends up being beneficial, the market allows the people who own the certificates to profit. But if the project ends up being harmful, the market does not inflict a cost on them. The certificates of a project that ended up being extremely harmful are worth as much as the certificates of a project that ended up being neutral, namely nothing. **Therefore, even if everyone believes that a certain project is net-negative, its certificates may be traded for a high price due to the chance that the project will end up being beneficial.**^[1]

Impact markets can thus incentivize people to create or fund net-negative projects. Denis Drescher [used](#)^o the term "distribution mismatch" to describe this risk, due to the mismatch between the probability distribution of investor profit and that of EV.

It seems especially important to prevent the risk from materializing in the domains of

anthropogenic x-risks and meta-EA. Many projects in those domains can cause a lot of accidental harm because, for example, they can draw attention to [info hazards](#), produce harmful outreach campaigns, produce dangerous experiments (e.g. in machine learning or virology), shorten AI timelines, intensify competition dynamics among AI labs, [etcetera](#).

Mitigating the risk is hard

The [Toward Impact Markets](#)^o post describes an approach that attempts to mitigate this risk. The core idea is that retro funders should consider the ex-ante EV rather than the ex-post EV if the former is smaller. (The details are more complicated; a naive implementation of this idea would incentivize people to launch a safe project and later expand it to include high-risk high-reward interventions.)

We think that this approach cannot be relied upon to sufficiently mitigate the risk due to the following reasons:

- For that approach to succeed, retro funders must be familiar with it and be sufficiently willing and able to adhere to it. However, some potential retro funders are more likely to use a much simpler approach, such as "you should buy impact that you like".
 - Other things being equal, simpler approaches are easier to communicate, more appealing to potential retro funders, more prone to become a meme and a norm, and more likely to be [advocated for](#)^o by teams who work on impact markets and want to get more traction.
- If there is no way to prevent anyone from becoming a retro funder, being careful about choosing/training the initial set of retro funders may not help much. Especially if the market allows people to profit from outreach interventions that attract new retro funders who are not very careful.
- The price of a certificate tracks the **maximum** amount of money that **any** future retro funder will be willing to pay for it. Prudent retro funders do not (significantly) offset the influence of imprudent retro funders on the prices of certificates of net-negative projects.
 - Traditional (prospective) charitable funding can have a similar dynamic; one only needs one funder to support a project even if everyone else thinks it's bad. Impact markets make the problem much worse, though, because they add

variance from uncertainty about project outcomes as well as variance in funder views.

- Suppose that a risky project that is ex-ante net-negative ends up being beneficial. If retro funders attempt to evaluate it after it already ended up being beneficial, [hindsight bias](#) can easily cause them to overestimate its ex-ante EV. This phenomenon can make the certificates of net-negative projects more appealing to investors, already at an early stage of the project (before it is known whether the project will end up being beneficial or harmful).

The conflict of interest problem for establishing impact markets

Can we just trust that people interested in establishing impact markets will do so only if it's a good idea? Unfortunately the incentivization of risky projects applies at this level. If someone establishes an impact market and it has large benefits, they might expect to be able to sell their impact in establishing the markets for large amounts of money. On the other hand if they establish impact markets and they cause large harms, they won't lose large amounts of money.

Establishing impact markets would probably involve many high-stakes decisions under great uncertainty. (e.g. should an impact market be launched? Should the impact market be decentralized? Should a certain person be invited to serve as a retro funder? Should certain certificates be deleted? What instructions should be communicated to potential market participants?) We should protect the integrity of these decisions by insulating them from conflicts of interest.

This point seems important even conditional on the people involved being the most careful and EA-aligned people in the world. (Because they are still humans, and humans' judgment is likely to be affected by biases/self-deception when there is a huge financial profit at stake).

Suggestions

- Currently, launching impact markets seems to us (non-robustly) net-negative. The

following types of impact markets seems especially concerning:

- Decentralized impact markets (in which there are no accountable decision makers that can control or shut down the market).
- Impact markets that allow certificates for risky interventions, and especially interventions that are related to the impact market itself (e.g. recruiting new retro funders).
- On the other hand, we're excited about work to further understand the benefits and costs of different funding structures. If there were a robust mechanism to allow the markets to avoid the risks discussed in this post (& ideally handle [moral trade](#)[◦] as well), we think impact markets could have very high potential. We just don't think we're there yet.
- In any case, **launching an impact market should not be done without (weak) consensus among the EA community**, in order to avoid the [unilateralist's curse](#)[◦].
- To avoid tricky conflicts of interest, **work to establish impact markets should only ever be funded in forward-looking ways**. Retro funders should commit to not buying impact of work that led to impact markets (at least work before the time when the incentivization of net-negative projects has been robustly cleared up, if it ever is). EA should socially disapprove of anyone who did work on impact markets trying to sell impact of that work.
- All of this relates to markets which encourage retrospective funding (especially but not exclusively if they also allow for the resale of impact).
 - In particular, this is not intended to apply to introducing market-like mechanisms like [explicit allocation of credit](#)[◦] between contributors to projects. While such mechanisms may be useful for supporting impact markets, they are also useful in their own right (for propagating price information without distorting incentives), and we're in favour of experiments with such credit allocation.

1. [^] The risk was probably first [pointed out](#) by Ryan Carey.



Mentioned in

- 112 **Impact Markets: The Annoying Details**
- 55 **Monthly Overload of EA - July 2022**
- 24 **A Fresh FAQ on GiveWiki and Impact Markets Generally**
- 22 **Crypto loves impact markets: Notes from Schelling Point Bogotá**
- 10 **Link Collection: Impact Markets**

[Load more \(5/6\)](#)

MORE POSTS LIKE THIS

- | | | | | |
|-----|---|---|-----------------|---|
| 54 | ▲ | Certificates of impact | Paul_Christiano | ⋮ |
| 85 | ▲ | Altruistic equity allocation | Paul_Christiano | ⋮ |
| 112 | ▲ | Impact Markets: The Annoying Details | Scott Alexander | ⋮ |

Comments 76

Sorted by [New & upvoted](#)

Some comments are truncated due to high volume. ([⌘F to expand all](#)) [⚙️ Change truncation settings](#)

▼ harsimony 3y [◀ 50 ▶](#) [✓ 0](#) [✗ 0](#) [😊](#)



I [proposed](#) a simple solution to the problem:

1. For a project to be considered for retroactive funding, participants must post a specific amount of money as collateral.
2. If a retroactive funder determines that the project was net-negative, they can burn the collateral to punish the people that participated in it. Otherwise, the project receives its collateral back.

This eliminates the "no downside" problem of retroactive funding and makes some net-negative projects unprofitable.

The amount of collateral can be chosen adaptively. Start with a small amount and increase it slowly until the number of net-negative projects is low enough. Note that setting the collateral too high can discourage net-positive but risky projects.

James Payor 3y < 18 > ✓ 0 ✗ 0 🗨️

Related: requiring some kind of insurance that pays out when a certificate becomes net-negative.

Suppose we somehow have accurate positive and negative valuations of certificates. We can have insurers sell put options on certificates, and be required to maintain that their portfolio has positive overall impact. (So an insurer needs to buy certificates of positive impact to offset negative impact they've taken on.)

Ultimately what's at stake for the insurer is probably some collateral they've put down, so it's a similar proposal.

mako yass 3y < 11 > ✓ 0 ✗ 0 🗨️

Crypto's inability to take debts or enact substantial punishments beyond slashing stakes is a huge limitation and I would like it if we didn't have to swallow that (ie, if we could just operate in the real world, with non-anonymous impact traders, who can be held accountable for more assets than they'd be willing to lock in a contract.)

Given enough of that, we would be able to implement this by just having an impact cert that's implicated in a catastrophe turn into debt/punishment, and we'd be able to make that disincentive a lot more proportional to the scale of its potential negative externalities, and we would be able to allow the market to figure out how big that risk is for itself, which is pretty much the point of an impact market.

Though, on reflection, I'm not sure I would want to let the market to decide that. The problem with markets is that they give us a max function, they're made of auctions, whoever pays most decides the price, and the views of everyone else are not taken into account at all. Markets, in a sense, subject us to the decisions of the people with the most extreme beliefs. Eventually the ones who are extreme and wrong go bankrupt and disappear, but I don't ... (read more)

7 Owen Cotton-Barratt 3y Nice, that's pretty interesting. (It's hacky, but that seems okay.) It's...

3 harsimony 3y This kind of thing could be made more sophisticated by making fines propor...

Stefan_Schubert 3y < 11 > ✓ 0 ✗ 0 🗨️

This kind of thing could be made more sophisticated by making fines proportional to the harm done

I was thinking of this. Small funders could then potentially buy insurance from large funders in

order to allow them to fund projects that they deem net positive even though there's a small risk of a fine that would be too costly for them.

▼ RyanCarey 3y < 22 > ✓ 0 ✗ 0 🗨️



I take it that Harsimony is proposing for the IC-seller to put up a flexible amount of collateral when they start their project, according to the possible harms.

There are two problems, though:

- This requires centralised prospective estimation of harms for every project. (A big part of the point of impact certificates is to evaluate things retroactively, and to outsource prospective evaluations to the market, thereby incentivising accuracy in the latter.
- This penalises IC-sellers based on how big their harms initially seem, rather than how big they eventually turn out to be.

It would be better if the IC-seller is required to buy insurance that will pay out the whole cost of the harm, as evaluated retrospectively. In order for the IC-seller to prove that they are willing to be accountable for all harms, they *must* buy insurance when they sell their IC. And to ensure that the insurer will pay out correctly, we must only allow insurers who use a standard, trusted board of longtermist evaluators to estimate the harms.

This means that a centralised system is only required to provide occasional retrospective evaluations of harm. The task of evaluating harms in prospect is delegated to insurers, similar to the role insurers play in the real world.

(This is my analysis, but the insurance idea was from Stefan.)

▼ RyanCarey 3y < 12 > ✓ 0 ✗ 0 🗨️



Although, the costs of insurance would need to be priced according to the ex ante costs, not the ex post costs.

For example: Bob embarks on a project with a 50% chance of success. If it succeeds, it saves one person's life, and Bob sells the IC. If it fails, it kills two people.

Clearly, the insurance needs to be priced to take into account a 50% chance of two deaths. So we would have to require Bob to buy the insurance when he initially embarks on the project (which is a tough ask, given that few currently anticipate selling their impact). Or else we would need to rely on a (centralised) retrospective evaluation of ex ante harm, for every project (which seems laborious).

8 Dawn Drescher 3y I love the insurance idea because compared to our previous ideas a...

4 Ofer 3y I don't think that short selling would work. Suppose a net-negative project has a ...

- 4 Emrik 3y I don't think such a rule has a chance of surviving if impact markets take off? 1. Ad...
- 4 mako yass 3y Traders would adopt a competitor without negative externality mechanisms,...
- 1 Emrik 3y But afaik the theory of change of this project doesn't rely on altruistic "end buye...
- 8 Linch 3y My understanding is that without altruistic end-buyers, then the intrinsic value ...
- 4 mako yass 3y There might be a market for that sort of ultimately valueless token now (o...
- 1 Emrik 3y Aye, I updated. I was kinda dumb. The magical speculation model is probably ...
- 2 mako yass 3y I think there's an argument for the thing you were saying, though... So...

▼ Dawn Drescher 3y < 30 > ✓ 0 ✗ 0 🗨️



Going Forward

1. We will convene a regular working group to more proactively iterate and improve the mechanism design focused on risk mitigation. We intend for this group to function for the foreseeable future. Anyone is welcome to join this group via our [Discord](#).
2. We will attempt to gain consultation from community figures that have expressed interest in impact markets (Paul Christiano, Robin Hanson, Scott Alexander, Eliezer Yudkowsky, Vitalik Buterin). This should move the needle towards more community consensus.
3. We will continue our current [EA Forum contest](#)^o. We will not run another contest in July.
4. We will do more outreach to other projects interested in this space (Bitcoin, Protocol Labs, Optimism, etc.) to make sure they are aware of these issues as well and we can come up with solutions together.

Do we think that impact markets are net-negative?

We – the Impact Markets team of Denis, Dony, and Matt – have been active EAs for almost a combined 20 years. In the past years we've individually gone through a prioritization process in which we've weighed importance, tractability, neglectedness, and personal fit for various projects that are close to the work of QURI, CLR, ACE, REG, CE, and o... (read more)

▼ Jan_Kulveit 3y < 21 > ✓ 0 ✗ 0 🗨️



If the main problem you want to solve is "scaling up grantmaking", there are probably many other ways how to do it other than "impact markets".

(Roughly, you can amplify any "expert panel of judges" evaluations with judgemental forecasting.)

- 3 Dawn Drescher 3y We've considered a wide range of mechanisms and ended up most opt...

- 1 Ofer 3y Does this mean that you (the Impact Markets team) may sell certificates of your wor...
- 1 Ofer 3y I think the analogy would work better if self-driving cars did risky things that could c...
- 3 Dawn Drescher 3y First of all, what we've summarized as "curation" so far could really be ...
- 1 Ofer 3y An impact market with invite-only access for issuers and investors seems safer t...

▼ Austin 3y < 25 > ✓ 0 ✗ 0 🗨️



Hm, naively - is this any different than the risks of net-negative projects in the for-profit startup funding markets? If not, I don't think this a unique reason to avoid impact markets.

My very rough guess is that impact markets should be at a bare minimum better than the for-profit landscape, which already makes it a worthwhile intervention. People participating as final buyers of impact will at least be looking to do good rather than generate additional profits; it would be very surprising to me if the net impact of that was worse than "the thing that happens in regular markets already".

Additionally - I think the negative externalities may be addressed with additional impact projects, further funded through other impact markets?

Finally: on a meta level, the amount of risk you're willing to spend on trying new funding mechanisms with potential downsides should basically be proportional to the amount of risk you see in our society at the moment. Basically, if you think existing funding mechanisms are doing a good job, and we're likely to get through the hinge of history safely, then new mechanisms are to be avoided and we want to stay the course. (That's not my current read of our xrisk situation, but would love to be convinced otherwise!)

- 7 Owen Cotton-Barratt 3y I think startups are usually doing an activity which scales if it's goo...
- 5 Stefan_Schubert 3y They refer to Drescher's post. He writes:
- 4 Owen Cotton-Barratt 3y I think this is not quite right. It shouldn't be about what we think ab...
- 2 Austin 3y I think impact markets should be viewed in that experimental lens, for what it's ...
- 3 Ofer 3y Impact markets can incentivize/fund net-negative projects that are not currently of ...

▼ Austin 3y < 12 > ✓ 0 ✗ 0 🗨️



Thanks for your responses!

I'm not sure that "uniqueness" is the right thing to look at.

Mostly, I meant: the for-profit world already incentivizes people to take high amounts of risk for financial gain. In addition, there are no special mechanisms to prevent for-profit entities from producing large net-negative harms. So asking that some special mechanism be introduced for impact-focused entities is an isolated demand for rigor.

There are mechanisms like pollution regulation, labor laws, etc which apply to for-profit entities - but these would apply equally to impact-focused entities too.

We should be cautious about pushing the world (and EA especially) further towards the "big things happen due to individuals following their local financial incentives" dynamics.

I think I disagree with this? I think people following local financial incentives is always going to happen, and the point of an impact market is to structure financial incentives to be aligned with what the EA community broadly thinks is good.

Agree that xrisk/catastrophe can happen via eg AI researchers following local financial incentives to make a lot of money - but unless your proposal is to overhaul the capitalist market system somehow, I think building a better competing alternative is the correct path forward.

8 Ofer 3y It may be useful to think about it this way: Suppose an impact market is launched...

1 Ofer 3y I messed up when writing that comment (see the EDIT block).

3 Owen Cotton-Barratt 3y I didn't follow this; could you elaborate? (/give an example?)

▼ vaniver 3y < 18 > ✓ 0 ✗ 0 🗨️



This reminds me a lot of limited liability (see also [Austin's comment°](#), where he compares it to the for-profit startup market, which because of limited liability for corporations bounds prices below by 0).

This is a historically unusual policy (full liability came first), and seems to me to have basically the same downsides (people do risky things, profiting if they win and walking away if they lose), and basically the same upsides (according to the theory supporting LLCs, there's *too little* investment and support of novel projects).

Can you say more about why you think this consideration is sufficient to be net negative? (I notice your post seems very 'do-no-harm' to me instead of 'here are the positive and negative effects, and we think the negative effects are larger', I'm also interested in Owen's impression on whether or not impact markets lead to more or less [phase 2 work°](#).)

1 Ofer 3y Can you explain the "same upsides" part? I think that most interventions that have a...

5 vaniver 3y Yeah; by default people have entangled assets which will be put at risk by starti...

2 Ofer 3y Efficient impact markets would allow anyone to create certificates for a project a...

▼ Yonatan Cale 3y < 16 > ✓ 0 ✗ 0 🗨️



Concretely:

If someone thinks a net-negative project is being traded on (or run at all), how about posting about it on the forum?

I assume anyone who retro funds a project will first run a search here and see what came up.

Meta:

1. The problem of funding net-negative projects exists also now.
2. I think that if retroactive public good's funding existed, people would be slightly horrified about canceling it. The upside seems big.
3. Nobody has seriously run retroactive public good's funding yet, we're not sure how it will go or what will go wrong. Surely some things will go wrong, I agree about that.
4. This is a kind of project that we can stop or change if we want to. There is a lot of human discretion. This is not like adding a government regulation that will be very hard to change, or launching a blockchain that you can pretty much never take back no matter what you do.
5. So I suggest we do experiment with it, fund some good projects (like hpmorpodcast!), and do a few things like encourage people to post on the forum if they think a project is net-negative and being traded. Perhaps even the software system that manages the trade can link to a forum post for each project, so people can voice their opinions there

- 9 Ofer 3y As we wrote in the post, even if everyone believes that a certain project is net-nega...
- 7 Yonatan Cale 3y OpenAI I recommend that early on someone posts: "OpenAI is expected t...
- 1 Ofer 3y But an impact market can still make OpenAI's certificates be worth \$100M if, for e...
- 3 Dawn Drescher 3y It could be done a bit more smoothly by (1) accepting no new issues, (2...
- 1 Ofer 3y That could make it easier for another team to create a new impact market that will...
- 2 Dawn Drescher 3y Okay, but to keep the two points separate: 1. Allowing people to mak...
- 1 Ofer 3y I think that a decentralized impact market that can't be controlled or shut down...

▼ mako yass 3y < 15 > ✓ 0 ✗ 0 🗨️



I should mention that Good Exchange/[impact certs](#) people have discussed this quite a bit. I raised concerns about this issue early on [here](#)°. Shortly later, I posted the question [would \(myopic\) general public good producers significantly accelerate the development of AGI?](#) to Lesswrong.

My current thoughts are similar to harsimony's, it's probably possible to get the potential negative externalities of a job to factor into the price of the impact cert by having certs take on negative value/turn into liabilities/debts if the negative outcomes end up eventuating.

We don't know exactly how to implement that well yet, though.

▼ Jordan Arel 3y < 8 > ✓ 0 ✗ 0 🗨️

I think that retroactive Impact Markets may be a net negative for many x-risk projects, however, I also think that in general Impact Markets may significantly reduce x-risk.

I think you have to bear in mind that if this project is highly successful, it has the potential to create a revolution in the funding of public goods. If humanity achieves much better funding and incentive mechanisms for public goods, this could create a massive increase in the efficiency of philanthropy.

It is hard to know how beneficial such a system would be, but it is not hard to s... (read more)

▼ Dawn Drescher 3y < 5 > ✓ 0 ✗ 0 🗨️

Dawn's (Denis's) Intellectual Turing Test Red-Teaming Impact Markets

[Edit: Before you read this, note that I failed. See the comment below.]

I want to check how well I understand Ofer's position against impact markets. The "Imagined Ofer" below is how I imagine Ofer to respond (minus language – I'm not trying to imitate his writing style though our styles seem similar to me). I would like to ask the real Ofer to correct me wherever I'm misunderstanding his true position.

I currently favor using the language of prize contests to explain impact markets unless ... (read more)

1 Ofer 3y I do not endorse the text written by "Imagined Ofer" here. Rather than describing all ...

▼ Jan_Kulveit 3y < 5 > ✓ 0 ✗ 0 🗨️

Seems a case of "[How x-risk projects are different from startups](#)"

▼ DC 3y < 4 > ✓ 0 ✗ 0 🗨️

This [post](#) seems like a really great example in the wild of how distribution mismatch could occur.

▼ Yonatan Cale 3y < 3 > ✓ 0 ✗ 0 🗨️

Another comment, specifically about AGI capabilities:

If someone wants to advance AI capabilities, they can already get prospective funding by opening a regular for-profit startup.

No?

1 Offer 3y Right. But without an impact market it can be impossible to profit from, say, publishi...

5 Yonatan Cale 3y I acknowledge this could be bad, but (as with most of my comments here...

3 Offer 3y It's just an example for how a post on the alignment forum can be net-negative a...

5 Yonatan Cale 3y I hope I'm not strawmanning your claim and please call me out if I am, ...

9 Offer 3y (Not an important point [EDIT: meaning the text you are reading in these paren...

2 Yonatan Cale 3y You changed my mind! I think the missing part, for me, is a public po...

1 DC 3y I think you're missing the part where if such a marketplace was materially chan...

▼ DC 3y < 3 > ✓ 0 ✗ 0 🗨️



Ofer (and Owen), I want to understand and summarize your cruxes one by one, in order to sufficiently pass your Ideological Turing Test that I can regenerate the core of your perspective. Consider me your point person for communications.

Crux: Distribution Mismatch of Impact Markets & Anthropogenic X-Risk

If I understand one of the biggest planks of your perspective correctly, you believe that there is a high-variance normal distribution of utility centered around 0 for x-risk projects, such that x-risk projects can often increase x-risk rather than decre... (read more)

3 Offer 3y I think that most interventions that have a substantial chance to prevent an existi...

▼ MichaelStJules 3y < 3 > ✓ 0 ✗ 0 🗨️



Based on your arguments here, I wouldn't be excited about impact markets. Still, if impact markets significantly expand the amount of work done for social good, it's plausible the additional good outweighs the additional bad. Furthermore, people looking to make money are already funding net negative companies due to essentially the same problems (companies have non-negative evaluations), so shifting them towards impact markets could be good, if impact markets have better projects than existing markets on average.

3 Offer 3y See my reply to Austin.

▼ DC 3y < 2 > ✓ 0 ✗ 0 🗨️



I am going to be extremely busy over the next week as I prep for the EAEcon retreat and wrap up the retro funding contest among other things, and combining that with the generally high personal emotional cost of engaging with this post will choose to not comment further for at least a week to focus my energy elsewhere (unless inspiration strikes me).

Here are a couple of considerations relevant to why I at least have not been more responsive, generally speaking:

- <https://forum.effectivealtruism.org/posts/oxZ5yaZKkwEswbCNs/limits-to-legibility°>
- <https://forum.eff°>

... (read more)

▼ Arepo 3y < 1 > ✓ 0 ✗ 0 😊⁺



Is there any real-world evidence of the unilateralist's curse being realised? My sense historically is that this sort of reasoning to date has been almost entirely hypothetical, and has done a lot to stifle innovation and exploration in the EA space.

7 Ofer 3y If COVID-19 is a result of a lab leak that occurred while conducting a certain type of ...

3 Arepo 3y I'm talking about the unilateralist's curse with respect to actions intended to be a...

4 Ofer 3y I don't find it hard to believe at all. Conditional on a lab leak, I'm pretty confident ...

3 Arepo 3y Strong disagree. A bioweapons lab working in secret on gain of function resea...

4 Ofer 3y I think that it's more likely to be the result of an effort to mitigate potential har...

2 Linch 3y I don't understand your sentence/reasoning here. Naively this should strengt...

2 Arepo 3y Why? The less scrupulous one finds Anthropic in their reasoning, the less ...

6 Linch 3y In 2015, when I was pretty new to EA, I talked to a billionaire founder of a company...

5 Ofer 3y Unless ~several people in EA had an opportunity to talk to that billionaire, I don't th...

▼ Linch 3y < 10 > ✓ 0 ✗ 0 😊⁺



Fair, though many EAs are probably in positions where they can talk to other billionaires (especially with >5 hours of planning), and probably chose not to do so.

More from Ofer

- ▲

73

Do AI companies make their safety researchers sign a non-dispara... Q

Ofer · 3y ago · 1m read

3

⋮
- ▲

63

If FTX is liquidated, who ends up controlling Anthropic? Q

Ofer · 3y ago · 1m read

8

⋮


View more

Curated and popular this week

Do primitive sentient organisms feel extreme pain?
disentangling intensity range and resolution

2 authors · 8d ago · 13m read · 9

Notes The following text explores, in a speculative manner, the evolutionary question: Did high-intensity affective states, specifically Pain, emerge early in evolutionary history, or did they develop gradually over time? Note: We are not neuroscientists; our work draws on our evolutionary biology background and our efforts to develop welfare metrics that accurately reflect reality and effectively reduce suffering. We hope these ideas may interest researchers in neuroscience, comparative cognition, and animal welfare science. This discussion is part of a broader ongoing project in progress, focusing on interspecific comparisons of affective capacities—a critical question for advancing Article 5 of the 1948 Universal Declaration of Human Rights states: "Obviously, no one shall be subjected to torture or to cruel, inhuman or degrading treatment or punishment." OK, it doesn't actually start with organisms experience extreme pain intensities, or fine-grained pain "obviously," but I like to imagine the commissioners all murmuring to intensity discrimination, or both? Scientific framing: Pain functions as a biological signalling system that guides behavior by encoding what the causal effect of Article 5 (or the 1984 UN Convention Against Torture) has been on reducing torture globally (though the physical range and resolution (i.e. the granularity with which differences in Pain intensity can be perceived) — can be viewed as an optimization problem, where neural architectures must balance computational efficiency, survival-driven signal prioritization, and adaptive flexibility (in another thread), David Thorstad wrote a mathematical clarification of the idea that small reductions in extinction risk have very high value, we haven't made much progress in the past 25 years. Reducing government torture and killings seems to be low in tractability. Despite reflecting also in granularity how finely intensity levels can be mathematics/. The arguments are many countries having a physical integrity rights index close to 1.0 (i.e.,



15 of 16

5/25/25, 5:48 PM

Recent opportunities in Building effective altruism

- ▲

11

EA Organization Updates: May 2025

Dane Valerie, Toby Tremlett ♦ · 11d ago · 9m read

💬

0

⋮
- ▲

15

Apply by 10th June: 'Introduction to Biosecurity' Online Course Starti...

Lin BL · 10d ago · 1m read

💬

0

⋮
- ▲

25

Twin Cities 10% Launch Event: Lives on the Line — Rapid Response F...

Ryan Begley ♦ · 7d ago · 1m read

💬

0

⋮

View more