

Rational Altruist

Adventures of a would-be do-gooder.

Why might the future be good?

by paulfchristiano

When talking about the future, I often encounter two (quite different) stories describing why the future might be good:

1. Decisions will be made by people whose lives are morally valuable and who want the best for themselves. They will bargain amongst each other and create a world that is good to live in. Because my values are roughly aligned with their aggregate preferences, I expect them to create a rich and valuable world (by my lights as well as theirs).
2. Some people in the future will have altruistic values broadly similar to my own, and will use their influence to create a rich and valuable world (by my lights as well as theirs).

Which of these pictures we take more seriously has implications for what we should do today. I often have object level disagreements which seem to boil down to disagreement about which of these pictures is more important, but rarely do I see serious discussion of that question. (When there is discussion, it seems to turn into a contest of political ideologies rather than facts.)

If we take picture (1) seriously, we may be interested in ensuring that society continues to function smoothly, that people are aware of and pursue what really makes them happy, that governments are effective, markets are efficient, externalities are successfully managed, etc. If we take picture (2) seriously, we are more likely to be concerned with changing what the people of the future value, bolstering the influence of people who share our values, and ensuring that altruists are equipped to embark on their projects successfully.

I'm mostly concerned with the very long run—I am wondering what conditions will prevail for most of the people who live in the future, and I expect most of them to be alive very far from now.

It seems to me that there are two major factors that control the relative importance of pictures (1) and (2): how prominent should we expect altruism to be in the future, and how efficiently are altruistic vs. selfish resources being used to create value? My answer to the second question is mostly vague hand-waving, but I think I have something interesting to say on the first question.

How much altruism do we expect?

I often hear people talking about the future, and the present for that matter, as if we are falling towards a Darwinian attractor of cutthroat competition and vanishing empathy (at least as a default presumption, which might be averted by an extraordinary effort). I think this picture is essentially mistaken, and my median expectation is that the future is much *more* altruistic than the present.

Dose natural selection select for self-interest?

In the world of today, it may seem that humans are essentially driven by self-interest, that this self-interest was a necessary product of evolution, that good deeds are principally pursued instrumentally in service of self-interest, and that altruism only exists at all because it is too hard for humans to maintain a believable sociopathic facade.

If we take this situation and project it towards a future in which evolution has had more time to run its course, creating automations and organizations less and less constrained by folk morality, we may anticipate an outcome in which natural selection has stripped away all empathy in favor of self-interest and effective manipulation. Some may view this outcome as unfortunate but inevitable, others may view it as a catastrophe which we should work to avert, and still others might view it as a positive outcome in which individuals are free to bargain amongst themselves and create a world which serves their collective interest.

But evolution itself does not actually seem to favor self-interest at all. *No matter what your values*, if you care about the future you are incentivized to survive, to acquire resources for yourself and your descendants, to defend yourself from predation, etc. etc. If I care about filling the universe with happy people and you care about filling the universe with copies of yourself, I'm not going to set out by trying to make people happy while allowing you and your descendants to expand throughout the universe unchecked. Instead, I will pursue a similar strategy of resource acquisition (or coordinate with others to stop your expansion), to ensure that I maintain a reasonable share of the available resources which I can eventually spend to help shape a world I consider value. (See [here](http://reflectivedisequilibrium.blogspot.com/2012/09/spreading-happiness-to-stars-seems.html) (<http://reflectivedisequilibrium.blogspot.com/2012/09/spreading-happiness-to-stars-seems.html>) for a similar discussion.)

This doesn't seem to match up with what we've seen historically, so if I claim that it's relevant to the future I have some explaining to do.

Historical distortions

Short-range consequentialism

One reason we haven't seen this phenomenon historically is that animals don't actually make decisions by backwards-chaining from a desired outcome. When animals (including humans) engage in goal-oriented behavior, it tends to be pretty local, without concern for consequences which are distant in time or space. To the extent that animal behavior is goal-oriented at a large scale, those goals are largely an emergent property of an interacting network of drives, heuristics, etc. So we should expect animals to have goals which lead them to multiply and acquire resources, even when those drives are pursued short-sightedly. And indeed, that's what we see. But it's not the fault of evolution alone—it is a product of evolution *given* nature's inability to create consequentialist reasoners.

Casually, we seem to observe a similar situation with respect to human organizations—organizations which value expansion for its own sake (or one of its immediate consequences) are able to expand aggressively, while organizations which don't value expansion have a much harder time deciding to expand for instrumental reasons without compromising their values.

Hopefully, this situation is exceptional in history. If humans ever manage to build systems which are properly consequentialist—organizations or automations which are capable of expanding because it is *instrumentally* useful—we should not expect natural selection to discriminate at all on the basis of those systems' values.

Value drift

Human's values are also distorted by the process of reproduction. A perfect consequentialist would prefer to have descendants who share their values. (Even if I value diversity or freedom of choice, I would like my children to at least share those values, at least if I want that freedom and diversity to last more than one generation!) But humans don't have this option—the only way we can expand our influence is by creating very lossy copies. And so each generation is populated by a fresh batch of humans with a fresh set of values, and the values of our ancestors only have an extremely indirect effect on the world of today.

Again, a similar problem afflicts human organizations. If I create a foundation that I would like to persist for generations, the only way it can expand its influence is by hiring new staff. And since those staff have a strong influence over what my foundation will do, the implicit values of my foundation will slowly but surely be pulled back to the values of the pool of human employees that I have to draw from.

These constraints distort evolution, causing selection to act only those traits which can be reliably passed on from one generation to the next. In particular, this exacerbates the problem from the preceding section—even to the extent that humans *can* engage in goal-oriented reasoning and expand their own influence instrumentally, these tendencies can not be very well encoded in genes or passed on to the next generation in other ways. This is perhaps the most fundamental change which would result from the development of machine intelligences. If it were possible to directly control the characteristics and values of the next generation, evolution would be able to act on those characteristics and values directly.

So what *does* natural selection select for?

If the next generation is created by the current generation, guided by the current generation's values, then the properties of the next generation will be disproportionately affected by those who care most strongly about the future.

In finance: if investors have different time preferences, those who are more patient will make higher returns and eventually accumulate much wealth. In demographics: if some people care more about the future, they may have more kids as a way to influence it, and therefore be overrepresented in future generations. In government: if some people care about what government looks like in 100 years, they will use their political influence to shape what the government looks like in 100 years rather than trying to win victories today.

What natural selection selects for is patience. In a thousand years, given efficient natural selection, the most influential people will be those who today cared what happens in a thousand years. Preferences about what happens to *me* (at least for a narrow conception of personal identity) will eventually die off, dominated by preferences about what society looks like on the longest timescales.

I think this picture is reasonably robust. There are ways that natural selection (/ efficient markets) can be frustrated, and I would not be *too* surprised if these frustrations persisted indefinitely, but nevertheless this dynamic seems like one of the most solid features of an uncertain future.

What values are we starting with?

Most of people's preferences today seem to concern what happens to them in the near term. If we take the above picture seriously, these values will eventually have little influence over society. Then the question becomes: if we focus only on humanity's collective preferences *over the long term*, what do those preferences look like? (Trying to characterize preferences as "altruistic" or not no longer seems useful as we zoom in.)

This is an empirical question, which I am not very well-equipped to value. But I can make a few observations that ring true to me (though my data is mostly drawn from academics and intellectuals, who may fail to be representative of normal people in important ways even after conditioning on the "forward-looking" part of people's values):

1. When people think about the far future (and thus when they articulate their preferences for the far future) they seem to engage a different mode of reasoning, more strongly optimized to produce socially praise-worthy (and thus prosocial) judgments. This might be characterized as a bias (<http://www.overcomingbias.com/2013/02/beware-far-values.html>), but to the extent we can talk about human preferences at all they seem to be a result of these kinds of processes (and to the extent that I am using my own altruistic values to judge futures, they are produced by a similar process). This effect seems to persist even when we are not directly accountable for our actions.
2. People mostly endorse their own enlightened preferences, and look discouragingly at attempts to

lock-in hastily considered values (though they often seem to have overconfident views about what their enlightened preferences will look like, which admittedly might interfere with their attempts at reflection).

3. I find myself sympathetic to very many people's accounts of their own preferences about the future, even where those accounts differ significantly from my own. I would be surprised if the distribution of moral preferences was too scattered.
4. To the extent that people care especially about their species, their nation, their family, themselves, etc.: they seem to be sensitive to fairness considerations (and rarely wish e.g. to spend a significant fraction of civilization's resources on themselves), their preferences seem to be only a modest distortion of aggregative values (wanting people with property X to flourish is not so different from wanting people to flourish, if property X is some random characteristic without moral significance), and human preferences seem to somewhat reliably drift in the direction of more universal concern as basic needs are addressed and more considerations are considered.

After cutting away all near-term interests, I expect that contemporary human society's collective preferences are similar to their stated moral preferences, with significant disagreement on many moral judgments. However, I expect that these values support reflection, that upon reflection the distribution of values is not too broad, and that for the most part these values are reasonably well-aligned. With successful bargaining, I expect a mixture of humanity's long-term interests to be only modestly (perhaps a factor of 10, probably not a factor of 1000) worse than my own values (as judged by my own values).

Moreover, I have strong intuitions to emphasize those parts of my values which are *least* historically contingent. (I accept that all of my values are contingent, but am happier to accept those values that are contingent on my biological identity than those that are contingent on my experiences as a child, and happier to accept those that are contingent on my experiences as a child than those that are contingent on my current blood sugar.) And I have strong reciprocity intuitions that exacerbate this effect and lead me to be more supportive of my peers' values. These effects make me more optimistic about a world determined by humanity's aggregate preferences than I otherwise would be.

How important is altruism?

(The answer to this question, unlike the first one, depends on your values: how important to what? I will answer from my own perspective. I have roughly aggregative values, and think that the goodness of a world with twice as many happy people is twice as high.)

Even if we know a society's collective preferences, it is not obvious what their relative *importance* is. At what level of prevalence would the contributions of explicit altruism become the source of value? If altruists are 10% of the influence-weighted population, do the contributions of the altruists matter? What if altruists are 1% of the population? A priori, it seems clear that the explicit altruists should do *at least* as much good—on the altruistic account—as any other population (otherwise they could decide to jump ship and become objectivists, or whatever). But beyond that, it isn't clear that altruists should create *much* more value—even on the altruistic account—than people with other values.

I suspect that explicit altruistic preferences create many times more value than self-interest or other

nearly orthogonal preferences. So in addition to expecting a future in which altruistic preferences play a very large role, I think that altruistic preferences would be responsible for most of the value even if they controlled only 1% of the resources.

One significant issue is population growth. Self-interest may lead people to create a world which is good for themselves, but it is unlikely to inspire people to create as many new people as they could, or use resources efficiently to support future generations. But it seems to me that the existence of large populations is a huge source of value. A barren universe is not a happy universe.

A second issue is that population *characteristics* may also be an important factor in the goodness of the world, and self-interest is unlikely to lead people to ensure that each new generation has the sorts of characteristics which would cause them to lead happy lives. It may happen by good fortune that the future is full of people who are well-positioned to live rich lives, but I don't see any particular reason this would happen. Instead, we might have a future "population" in which almost all resources support automation that doesn't experience anything, or a world full of minds which crave survival but experience no joy, or etc.; "self-interest" wouldn't lead any of these populations to change themselves to experience more happiness. It's not clear why we would avoid these outcomes except by a law of nature that said that productive people were happy people (which seems implausible to me) or by coordinating to avoid these outcomes.

(If you have different values, such that there *is* a law [or at least guideline] of nature: "productive people are morally valuable people," then this analysis may not apply to you. I know several such people, but I have a hard time sympathizing with their ethics.)

Conclusion

I think that the goodness of a world is mostly driven by the amount of explicit optimization that is going on to try and make the world good (this is all relative to my values, though a similar analysis seems to carry with respect to other aggregative values). This seems to be true even if relatively little optimization is going on. Fortunately, I also think that the future will be characterized by *much* higher influence for altruistic values. If I thought altruism was unlikely to win out, I would be concerned with changing that. As it is, I am instead more concerned with ensuring that the future proceeds without disruptions. (Though I still think it is worth it to try and increase the prevalence of altruism faster, most of all because this seems like a good approach to minimizing the probability of undesired disruptions.)

PUBLISHED: [February 27, 2013 \(2013-02-27T05:15:05+0000\)](#)

FILED UNDER: Looking ahead

11 Comments to "Why might the future be good?"

Peli Grietzer says:

March 9, 2013 at 2:48 am

I'm not sure why you suggest that reflection reduces the divergence in people's long-term

preferences. My impression has always been that reflection causes radical divergence in people's long-term preferences, while leaving their short-term preferences roughly as converged as before. Two attempted examples:

If we include both people in the rationalist community and philosophically informed people outside the community as examples of people who have reflected about their long-term preferences, then reflection seems to lead to a strong divergence in population ethics. Almost every philosophically informed person I know outside the rationalist community accepts some form of asymmetry thesis to the effect that extra good lives aren't in themselves a major improvement of the world, whereas no one I know of in the rationalist community does.

If we include just the rationalist community, then there's more agreement about population ethics (converging on something roughly aggregative), but there's strong divergence about what makes lives good. Hedonism seems almost-dominant with people around the FHI, desire-satisfactionism seems dominant with people around OB, and pluralism seems dominant with people around LW. Reflection seems to have split people in the rationalist community into hedonists, desire-satisfactionists, and pluralists, and it is fairly hard to come up with a future that none of the three will consider astronomically wasteful.

REPLY

Ruairí Donnelly says:

April 12, 2013 at 6:19 pm

If you consider non-human animals too (and there are great reasons to do so, and basically no reasons not to: <http://vimeo.com/58451841>) then it's a much better idea to try and convince more people of your personal ethic. The values of society are very bad at the moment and it's likely humanity may create a very bad future, see more here:

<http://felicifia.org/viewtopic.php?p=4454>

and

<http://www.utilitarian-essays.com/values-spreading.html>

and

<http://www.utilitarian-essays.com/against-wishful-thinking.html>

REPLY

Brian Tomasik says:

May 7, 2013 at 3:14 am

Very interesting stuff, Paul! It's important to think about what kinds of broad trends we can anticipate based on the dynamics at play, without going into overly specific scenarios.

My main disagreement with your conclusions comes from the facts that (a) I don't expect a lot of convergence among human values and (b) to the extent there is convergence, I expect to be on the losing end of it, because I have significantly different values from most smart people I know.

When you have weird values, spreading those values is often going to be better than increasing/decreasing the chance of survival. You seem to mainly disagree with this because you're mostly ok with the range of values that you think might be produced. Still, you say the actual values of the

future could be ~10 times worse than your idealized ones. Isn't that a pretty significant margin? Spreading your values given human survival would be 10 times more important than ensuring human survival in general, to say nothing of the point that it's probably easier to spread your values a given amount than to ensure survival a given amount.

Finally, what do you make of the two concerns you mention in the "How Important is Altruism" section? Those both seem like appreciable worries for you. Might there be higher leverage in tackling those?

REPLY

Against moral advocacy | Rational Altruist says:

June 13, 2013 at 9:56 pm

[...] that because most people's self-interest seems to be relatively satiable and short-sighted, I expect the self-interested parts of people's motivations to mostly trade vast resources in the [...]

Reply

Ems Evolve | Bayesian Investor Blog says:

July 28, 2017 at 9:35 pm

[...] also Paul Christiano's Why might the future be good?. ("What natural selection selects for is [...]

Reply

Jan says:

November 9, 2017 at 10:13 am

Hey Paul, awesome post. I am a few years late to the party, but better late than never :D. There are a few things I don't quite understand about the post yet:

Why/how will machine intelligence make it possible to directly control the values and characteristics of the next generation? (this seems to be a necessary condition for your argument)

Also: Will your argument break down if immortality is invented? In this case, evolution would still select for patience, but selfish actors might have strong preferences about the long-term outcome and be patient.

REPLY

Cause prioritization for downside-focused value systems – Foundational Research Institute says:

February 21, 2018 at 12:12 pm

[...] some islands of happiness (or even have a sea of happiness and some islands of suffering). See also these posts on why the future is likely to be net good in expectation according to views where creating [...]

Reply

elityre says:

July 13, 2018 at 7:33 pm

Ok. I take your basic claim, in the first part of this essay to be something like...

"Currently, natural selection pressures do tend towards self-interest, but this is only because current organisms and organizations are currently implemented in such a way that their actions are heavily contingent on their values. Humans, in practice, are NOT mostly consequentialist agents for whom it is psychologically plausible to take whichever action has the highest expected value (according to

their own values).

Instead, human psychology is built so that the values and the behaviors are both emergent properties of a soup of heuristics, and as a result, humans don't efficiently take advantage of instrumental convergence. It isn't psychologically plausible for humans to pursue goals arbitrary instrumental goals with the fervor that they pursue things that seem valuable in of themselves.

A consequentialist architecture, dominates a "heuristic soup" architecture, so we should expect selection pressures for consequentialist agents with convergent instrumental goals. And once consequentialism reaches fixation, there's virtually no selection pressure on the values of agents, only on their rationality.

To put it another way, humans aren't rational, so what values they have strongly influences their behavior. We should expect the future to be dominated by rational agents, so their values will not impact their (short run) behavior much."

(Please correct me if my paraphrase is off in some way.)

This is all good and fine, but if I put on my Robin Hanson frame, I would say that this dynamic doesn't affect the way the future goes very much.

It seems like this model implies a future where consequentialist agents first scramble to capture as much of the resources of the universe as possible, then, after the universe has been carved up and is in equilibrium, everyone switches to exploit mode and uses their accumulated resources in accordance with their values. First a "career" of acquiring resources, then a "retirement" of implementing one's values.

I'm crafting a response, but I first want to confirm that I understand the basic picture.

REPLY

Reducing Risks of Astronomical Suffering: A Neglected Priority – Foundational Research Institute

says:

August 26, 2019 at 2:49 pm

[...] correct to assume that the ratio of expected future happiness to suffering is greater than one, and plausibly quite a bit larger than that,⁷ the case is not open-and-shut. Good values alone are not sufficient for ensuring good [...]

Reply

Reducing Risks of Astronomical Suffering: A Neglected Priority – Center on Long-Term Risk says:

March 1, 2020 at 1:59 pm

[...] correct to assume that the ratio of expected future happiness to suffering is greater than one, and plausibly quite a bit larger than that,⁷ the case is not open-and-shut. Good values alone are not sufficient for ensuring good [...]

Reply

Cause prioritization for downside-focused value systems – Center on Long-Term Risk says:

April 28, 2020 at 1:23 pm

[...] some islands of happiness (or even have a sea of happiness and some islands of suffering). See also these posts on why the future is likely to be net good in expectation according to views where creating [...]

Reply

Blog at WordPress.com. Do Not Sell or Share My Personal Information